

Learning with previously unseen features

Advisor- Naresh Manwani

Mentor-Kulin Shah

Team No.-11

Nakul Vaidya 201501108

Omkar Miniar 201530217

Ritvick Gupta 201530057

Akshay Vyas 201530032

References :

Learning with with Previously Unseen Features by YUAN SHI and CRAIG A. KNOBLOCK (IJCAI - 17):

<https://www.ijcai.org/proceedings/2017/0379.pdf>

GITHUB LINK :

https://github.com/akshay25vyas/learning_from_unseen_features

Problem Statement :

- **Improving** a machine learning model by identifying and using features that are **not in training set**.
- We propose a novel approach that learns a model over **both original** and **new features**.
- Author's result of LUF shows significant improvement over baselines.

What is different ?

- Our problem is a special case of **semi-supervised learning**.
- Existing semi-supervised learning algorithms focus on using unlabeled data from the same feature space of labeled data .
- But we are building a new machine learning model that models over both original and new-features .

Ex : Consider a model that predicts a job applicant's quality . Now qualities will be modeled using both qualities on resume and say applicant's social media such as facebook.

Approach :

Given :

N Labeled samples $\{(X_s, Y_s)\}_s^N$ and M Unlabeled samples $\{(X_t, Z_t)\}_t^M$

Predict:

We want to learn a model $f_\theta(x, z)$ where θ represents the model parameters, and the predicted label as $\hat{y} = f_\theta(\mathbf{X}, \mathbf{Z})$

Challenges :

- For each source-domain sample (X_s, Y_s) , if we could estimate its \hat{z}_s reliably, we can simply train $f_{\theta}(x,z)$ on the source domain. However, estimating z from x can be challenging when their dependency is weak.
- Training on the target domain is very challenging as there are no labels available.

Solutions:

- If our model predicts target-domain labels $\{\hat{y}_t\}$ well, then $\{\hat{y}_t\}$ should be consistent with the training labels. Consistency is expressed through the joint distribution of (x, y) such that $\{(x_s, y_s)\}$ and $\{(x_t, \hat{y}_t)\}$ are mixed as much as possible.
- When this happens, each source-domain sample (x_s, y_s) becomes close to its k -nearest neighbors in the target domain, and vice-versa. Therefore, we propose the following objective function to minimize the cross-domain k -nearest neighbor distances in the joint space of (x, y)

Problem Formulation :

$$\begin{aligned} \min_{\theta} \sum_s \sum_{t \in \mathcal{N}_{\mathcal{T}}^k(s)} \text{dist}[(\mathbf{x}_s, y_s), (\mathbf{x}_t, \hat{y}_t)] \\ + \sum_t \sum_{s \in \mathcal{N}_{\mathcal{S}}^k(t)} \text{dist}[(\mathbf{x}_t, \hat{y}_t), (\mathbf{x}_s, y_s)] + \lambda \|\theta\|_2^2 \end{aligned}$$

Where ,

$$\text{dist}[(\mathbf{x}_s, y_s), (\mathbf{x}_t, \hat{y}_t)] = \|\mathbf{x}_s - \mathbf{x}_t\|_2^2 + \gamma \Delta(y_s, \hat{y}_t)$$

Page _____

Regularized Linear Regression:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

Now,

$$\frac{\partial}{\partial \theta_j} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\Rightarrow 2 \left[(h_{\theta}(x^{(i)}) - y^{(i)}) \frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)}) \right]$$

For linear regression model

$$\frac{\partial}{\partial \theta_j} (h_{\theta}(x^{(i)})) = [x^{(i)}]_j$$

$$\frac{\partial}{\partial \theta_j} \lambda \sum_{j=1}^n \theta_j^2 = 2\lambda \theta_j$$

So for regularized linear case,

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \lambda \theta_j \right]$$

$$\begin{aligned} \min_{\theta} \sum_s \sum_{t \in \mathcal{N}_{\mathcal{T}}^k(s)} [v_{st}^2 + \gamma \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] \\ + \sum_t \sum_{s \in \mathcal{N}_{\mathcal{S}}^k(t)} [v_{ts}^2 + \gamma \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] + \lambda \|\theta\|_2^2 \end{aligned}$$

For regression tasks, we simply set :

$$\Delta(y_s, \hat{y}_t) = \|y_s - \hat{y}_t\|^2.$$

For classification tasks with C classes, we use probabilistic classification models and set :

$$\Delta(y_s, \hat{y}_t) = 1 - \sum_{c=1}^C y_s(c) \hat{y}_t(c).$$

Where \hat{y}_t is a C-dimensional vector representing the probability in each class, and y_s is a C-dimensional binary vector.

Alternating Optimization :

Let $V_{\mathcal{T}}^k(s)$ index (X_s, Y_s) 's any (not necessarily the nearest) k neighbors in the target domain, and $V_s^k(t)$ index (x_t, y_t) 's any k neighbors in the source domain.

$$\begin{aligned} & \sum_{t \in \mathcal{N}_{\mathcal{T}}^k(s)} [v_{st}^2 + \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] \\ &= \min_{\mathcal{V}_{\mathcal{T}}^k(s)} \sum_{t \in \mathcal{V}_{\mathcal{T}}^k(s)} [v_{st}^2 + \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] \end{aligned}$$

This is equivalent to

$$\begin{aligned} & \min_{\theta, \{\mathcal{V}_{\mathcal{T}}^k(s)\}, \{\mathcal{V}_s^k(t)\}} \sum_s \sum_{t \in \mathcal{V}_{\mathcal{T}}^k(s)} [v_{st}^2 + \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] \\ & + \sum_t \sum_{s \in \mathcal{V}_s^k(t)} [v_{ts}^2 + \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t))] + \lambda \|\theta\|_2^2 \end{aligned}$$

When θ is fixed, we update $\{V_{\mathcal{T}}^k(s)\}$ and $\{V_s^k(t)\}$ based on nearest neighbor search.
When $\{V_{\mathcal{T}}^k(s)\}$ and $\{V_s^k(t)\}$ are fixed, we optimize θ by solving

$$\begin{aligned} \min_{\theta} \sum_s \sum_{t \in \mathcal{V}_{\mathcal{T}}^k(s)} \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t)) \\ + \sum_t \sum_{s \in \mathcal{V}_{\mathcal{S}}^k(t)} \Delta(y_s, f_{\theta}(\mathbf{x}_t, \mathbf{z}_t)) + \lambda \|\theta\|_2^2 \end{aligned}$$

This is easier to optimize than earlier equation when f_{θ} is smooth in θ .

Algorithm 1 Optimization algorithm for LUF

Input: source-domain samples $\{(\mathbf{x}_s, y_s)\}_{s=1}^N$ and target-domain samples $\{(\mathbf{x}_t, \mathbf{z}_t)\}_{t=1}^M$, neighbor size k , weight parameter γ , and regularization parameter λ .

Initialize θ by solving Eq. (8)

for $iter = 1, 2, \dots, T$ **do**

for $s = 1, 2, \dots, N$ **do**

 Fix θ , update $\mathcal{V}_{\mathcal{T}}^k(s)$

for $t = 1, 2, \dots, M$ **do**

 Fix θ , update $\mathcal{V}_{\mathcal{S}}^k(t)$

 Fix $\{\mathcal{V}_{\mathcal{T}}^k(s)\}, \{\mathcal{V}_{\mathcal{S}}^k(t)\}$, optimize θ by solving Eq. (6)

Output: a local optimal solution θ^* .

Atom File Edit View Selection Find Packages Window Help

classification.py — ~/sem6/smai/finalproject

Project

- file1000.txt
- filetest.txt
- filetrain.txt
- test.txt
- train.txt
- usps
 - usps.jpg
 - usps.t
- .DS_Store
- bibtex-379.bib
- change_nakul.py
- change_omkar.py
- classification.py
- dataparser.py
- filetest.txt
- filetrain.txt
- house_files.zip
- KernelRidge.py
- meracode.py
- nakul.py
- naya.py
- pr.py
- proj.py
- project
- project.pdf
- regression.py
- temp1.py
- temp2.py

classification.py

```
181 yf2=np.zeros((k*m,noc),'float')
182
183 T=30 # no. of times you want to run iteration for updating thetha
184 for it in range(30):
185     #print it
186     #print "iteration"
187     for s in range(len(xs)):
188         V_t_k[s]=upda_nn_target(tomi,xs[s],ysf[s,:],xt,zt,gamma,k) # for every xs[s] update its nearest neighbours in target
189         for j in range(k):
190             yf1[s*k+j]=ysf[s]
191         #print "fjkhasdfjkhsd"
192         for t in range(len(xt)):
193             V_s_k[t]=upda_nn_source(tomi,xt[t],zt[t],xs,ysf,gamma,k) # for every xt[t] update its nearest neighbours in source do
194             for j in range(k):
195                 xf2[t*k+j,0:d]=xt[t]
196                 xf2[t*k+j,d]=zt[t]
197
198         xf1=np.reshape(V_t_k,(m*k,d+2))
199         yf2=np.reshape(V_s_k,(m*k,10))
200         finalx=np.vstack([xf1,xf2])
201         #print yf1.shape
202         #print yf2.shape
203         finaly=np.vstack([yf1,yf2])
204         #print finalx.shape
205         # now we have V_t_k & V_s_k ,update thetha
206         m1,d1=finalx.shape
207         yuu=np.identity(d1)
208         yuu[0,0]=0.0
209         on=np.ones((m1,1))
210         #finalx=np.hstack([finalx,on])
211         tomi=np.dot(np.dot(np.linalg.pinv(np.dot(np.transpose(finalx),finalx)+ld*yuu),np.transpose(finalx)),finaly)
212         #tomi=np.dot(np.dot(np.linalg.pinv(np.dot(np.transpose(np.hstack([finalx,on])),np.hstack([finalx,on]))+ld*yuu),np.transpo
213
214         y_pred=np.dot(finalx,tomi)
215         answer=np.sqrt(mean_squared_error(finally, y_pred))
216         tomi=np.transpose(tomi)
217         #print tomi.shape
218         print answer
219
```

classification.py 85:26

LF UTF-8 Python master 26 files 1 update

Atom File Edit View Selection Find Packages Window Help

regression.py — ~/sem6/smai/finalproject

Project

file1000.txt

filetest.txt

filetrain.txt

test.txt

train.txt

usps

usps.jpg

usps.t

.DS_Store

bibtex-379.bib

change_nakul.py

change_omkar.py

classification.py

dataparser.py

filetest.txt

filetrain.txt

house_files.zip

KernelRidge.py

meracode.py

nakul.py

naya.py

pr.py

proj.py

project

project.pdf

regression.py

temp1.py

temp2.py

classification.py

regression.py

```
145 | /Users/apple/sem6/smai/finalproject/regression.py
146 V_t_k=np.zeros((m,k,d+1),'float')
147 V_s_k=np.zeros((m,k,1),'float')
148
149 xf1=np.zeros((k*m,d+1),'float')
150 xf2=np.zeros((k*m,d+1),'float')
151 yf1=np.zeros((k*m,1),'float')
152 yf2=np.zeros((k*m,1),'float')
153 tomi=np.zeros((d+2,1),'float')
154 omi=np.zeros((d+2,1),'float')
155 regg=np.zeros((d+2,d+2),'float')
156 T=160 # no. of times you want to run iteration for updating thetha
157 for it in range(T):
158     for s in range(len(xs)):
159         V_t_k[s]=upda_nn_target(thetha,intercept,xs[s],ys[s],xt,zt,gamma,k) # for every xs[s] update its nearest neighbours in
160         for j in range(k):
161             yf1[s*k+j]=ys[s]
162     for t in range(len(xt)):
163         V_s_k[t]=upda_nn_source(thetha,intercept,xt[t],zt[t],xs,ys,gamma,k) # for every xt[t] update its nearest neighbours in
164         for j in range(k):
165             xf2[t*k+j,0:d]=xt[t]
166             xf2[t*k+j,d]=zt[t]
167
168     xf1=np.reshape(V_t_k,(m*k,d+1))
169     yf2=np.reshape(V_s_k,(m*k,1))
170     finalx=np.vstack([xf1,xf2])
171     finaly=np.vstack([yf1,yf2])
172     m1,d1=finalx.shape
173     yuu=np.identity(d1+1)
174     yuu[0,0]=0.0
175     on=np.ones((m1,1))
176     finalx=np.hstack([finalx,on])
177     tomi=np.dot(np.dot(np.linalg.pinv(np.dot(np.transpose(finalx),finalx)+ld*yyu),np.transpose(finalx)),finaly)
178     y_pred=np.dot(finalx,tomi)
179
180     answer=np.sqrt(mean_squared_error(finaly, y_pred))
181     omi=copy.deepcopy(tomi)
182
```

regression.py 145:5

LF UTF-8 Python master 26 files 1 update

Initialization :

The quality of the solution depends on how we initialize θ . If we could get a good estimate of \hat{z} 's based on each source-domain sample x_s , we can initialize θ by solving

$$\min_{\theta} \sum_s \Delta(y_s, f_{\theta}(x_s, \hat{z}_s)) \quad (7)$$

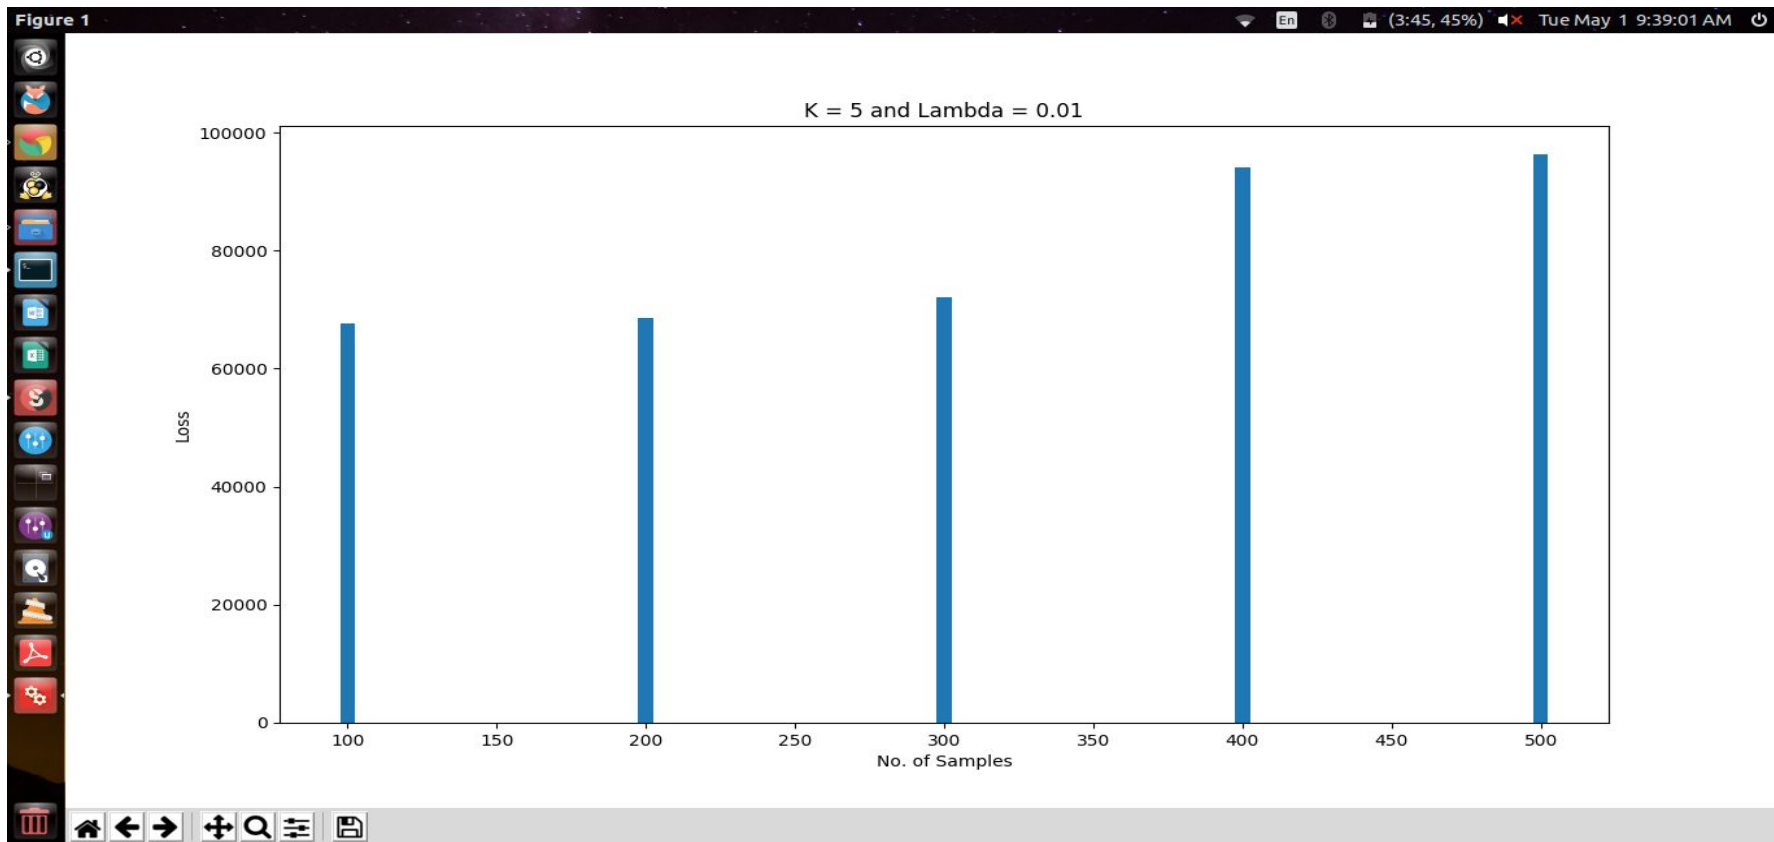
However, estimating \hat{z} 's can be very challenging when the dependency between x and z is weak. for each x_s , we find a set of its nearest neighbors in $\{x_t\}$, and use the corresponding z_t to form a candidate set Z_s . We then minimize the model error by optimizing both θ and $\{\hat{z}_s\}$:

$$\min_{\theta} \sum_s \min_{\hat{z}_s \in Z_s} \Delta(y_s, f_{\theta}(x_s, \hat{z}_s)) \quad (8)$$

Complexity Analysis:

- In our Algorithm, each iteration involves N updates on $V_T^k(s)$ and M updates on $V_s^k(t)$.
- Each update on $V_T^k(s)$ takes $O(MD^2)$ and each update on $V_s^k(t)$ takes $O(ND^2)$, where D is the dimensionality of original features. Therefore the complexity of updating $\{V_T^k(s)\}$ and $\{V_s^k(t)\}$ at each iteration is $O(NMD^2)$.
- Additionally, each iteration involves learning a linear regression or logistic regression function, whose complexity is $O((D + W)^2(N + M))$ where W is the dimensionality of new features and we assume $N > D + W$.
- Further assuming $W = O(D)$ and $M = O(N)$, we have the overall complexity as $O(D^2N^2)$.

Results:



Results:

For Regression phase:

% of increase as compared to diff model:-

We used kernel Regression: 164071.171593 was its rmse error.

Our model Rmse Error : 135234.050503 (for 500 samples)

% increase is : $(164071.171593 - 135234.050503) / 164071.171593 * 100$

I.e - 17.57% of improvement.

Results:

For Classification phase:

% of increase as compared to diff model:-

We used Logistic Regression: 13.256% was its classification error rate.

Our model classification error rate : 12.586% (for 500 samples)

% increase is : $(13.256 - 12.586) / 13.256 * 100$

I.e - 5.05% of improvement.

Work Done:

- Implemented complete Model with Alternating optimization.
- Implemented both regression and classification model.
- Dataset used for calculating prediction error (RMSE) on
 - Regression : House, 8 for housing price prediction, contains 20,640 samples with 9 features.
 - Classification: USPS, which recognizes handwriting digits from images, contains 9,298 samples from 10 classes.

Conclusion

- We presented a novel machine learning approach that leverages previously unseen features in the training set. The approach is applicable to both classification and regression tasks. Supported by our empirical results, the approach can be used to improve a learning model when new features are accessible.
- We also plan to develop algorithms that can automatically determine when to explore new features and which features to select from a large pool of features in an open environment.

Thank You

